

Class 17: Mini Project

Christina Mac

We will start by downloading the most recently dated “Statewide COVID-19 Vaccines Administered by ZIP code” CSV file:

```
# Import vaccination data
vax <- read.csv("covid19vaccinesbyzipcode_test.csv")
head(vax)
```

	as_of_date	zip_code_tabulation_area	local_health_jurisdiction	county
1	2021-01-05	95446	Sonoma	Sonoma
2	2021-01-05	96014	Siskiyou	Siskiyou
3	2021-01-05	96087	Shasta	Shasta
4	2021-01-05	96008	Shasta	Shasta
5	2021-01-05	95410	Mendocino	Mendocino
6	2021-01-05	95527	Trinity	Trinity

	vaccine_equity_metric_quartile	vem_source
1	2	Healthy Places Index Score
2	2	CDPH-Derived ZCTA Score
3	2	CDPH-Derived ZCTA Score
4	NA	No VEM Assigned
5	3	CDPH-Derived ZCTA Score
6	2	CDPH-Derived ZCTA Score

	age12_plus_population	age5_plus_population	tot_population
1	4840.7	5057	5168
2	135.0	135	135
3	513.9	544	544
4	1125.3	1164	NA
5	926.3	988	997
6	476.6	485	499

	persons_fully_vaccinated	persons_partially_vaccinated
1	NA	NA
2	NA	NA

3	NA	NA
4	NA	NA
5	NA	NA
6	NA	NA
percent_of_population_fully_vaccinated		
1	NA	
2	NA	
3	NA	
4	NA	
5	NA	
6	NA	
percent_of_population_partially_vaccinated		
1	NA	
2	NA	
3	NA	
4	NA	
5	NA	
6	NA	
percent_of_population_with_1_plus_dose		booster_recip_count
1	NA	NA
2	NA	NA
3	NA	NA
4	NA	NA
5	NA	NA
6	NA	NA
bivalent_dose_recip_count		eligible_recipient_count
1	NA	0
2	NA	0
3	NA	2
4	NA	2
5	NA	0
6	NA	0
redacted		
1	Information redacted in accordance with CA state privacy requirements	
2	Information redacted in accordance with CA state privacy requirements	
3	Information redacted in accordance with CA state privacy requirements	
4	Information redacted in accordance with CA state privacy requirements	
5	Information redacted in accordance with CA state privacy requirements	
6	Information redacted in accordance with CA state privacy requirements	

Q1. What column details the total number of people fully vaccinated? Column 11, “persons_fully_vaccinated”

Q2. What column details the Zip code tabulation area? zip_code_tabulation_area

Q3. What is the earliest date in this dataset? 2021-01-05

```
head(sort(vax$as_of_date))
```

```
[1] "2021-01-05" "2021-01-05" "2021-01-05" "2021-01-05" "2021-01-05"  
[6] "2021-01-05"
```

Q4. What is the latest date in this dataset? 2023-02-28

```
head(sort(vax$as_of_date, decreasing = TRUE))
```

```
[1] "2023-02-28" "2023-02-28" "2023-02-28" "2023-02-28" "2023-02-28"  
[6] "2023-02-28"
```

We can use the skim() function for a quick overview of a new dataset like this:

```
skimr::skim(vax)
```

Table 1: Data summary

Name	vax
Number of rows	199332
Number of columns	18
Column type frequency:	
character	5
numeric	13
Group variables	None

Variable type: character

skim_variable	n_missing	complete_rate	min	max	empty	n_unique	whitespace
as_of_date	0	1	10	10	0	113	0
local_health_jurisdiction	0	1	0	15	565	62	0
county	0	1	0	15	565	59	0
vem_source	0	1	15	26	0	3	0

skim_variable	n_missing	complete_rate	min	max	empty	n_unique	whitespace
redacted	0	1	2	69	0	2	0

Variable type: numeric

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
zip_code_tabulation_area	0	1.00	93665.11	1817.38	0	192257.75	3658.50	5380.50	7635.0	
vaccine_equity_metric_083tile	0	0.95	2.44	1.11	1	1.00	2.00	3.00	4.0	
age12_plus_population	0	1.00	18895.01	8993.87	0	1346.95	13685.13	1756.18	8556.7	
age5_plus_population	0	1.00	20875.24	1105.97	0	1460.50	15364.00	1877.00	1902.0	
tot_population	9718	0.95	23372.77	2628.51	2	2126.00	18714.08	168.00	1165.0	
persons_fully_vaccinated	16525	0.92	13962.35	5054.09	1	930.00	8566.00	23302.08	7566.0	
persons_partially_vaccinated	16525	0.92	1701.64	2030.18	11	165.00	1196.00	2535.00	39913.0	
percent_of_population_fully_vaccinated	16525	0.90	0.57	0.25	0	0.42	0.60	0.74	1.0	
percent_of_population_partially_vaccinated	16525	0.90	0.08	0.09	0	0.05	0.06	0.08	1.0	
percent_of_population_1_plus_dose	16525	0.89	0.63	0.24	0	0.49	0.67	0.81	1.0	
booster_recip_count	72872	0.63	5837.31	7165.81	11	297.00	2748.00	438.25	9553.0	
bivalent_dose_recip_count	158664	0.20	2924.93	3583.45	11	190.00	1418.00	1626.25	7458.0	
eligible_recipient_count	0	1.00	12801.84	4908.33	0	504.00	6338.00	21973.08	7234.0	

Q5. How many numeric columns are in this dataset? 13

Q6. Note that there are “missing values” in the dataset. How many NA values there in the persons_fully_vaccinated column? 16525

```
sum(is.na(vax$persons_fully_vaccinated))
```

```
[1] 16525
```

What percent of persons_fully_vaccinated values are missing (to 2 significant figures)? 8.3%

```
length(vax$persons_fully_vaccinated)
```

```
[1] 199332
```

```
round((sum(is.na(vax$persons_fully_vaccinated))/length(vax$persons_fully_vaccinated))*100,
```

```
[1] 8.29
```

```
##Working with dates
```

the lubridate package makes working with dates and times in R much less of a pain. Let's have a first play with this package here.

```
library(lubridate)
```

```
Attaching package: 'lubridate'
```

The following objects are masked from 'package:base':

```
date, intersect, setdiff, union
```

```
today()
```

```
[1] "2023-03-07"
```

```
# This will give an Error!  
#today() - vax$as_of_date[1]
```

```
# Specify that we are using the year-month-day format  
vax$as_of_date <- ymd(vax$as_of_date)
```

We can now magically do math with dates

```
today()-ymd("2021-01-05")
```

Time difference of 791 days

How old am I?

```
today()-ymd("2001-02-11")
```

Time difference of 8059 days

Let's treat the whole col

How many days have passed since the first vaccination reported in this dataset?

```
today()-ymd(vax$as_of_date[1])
```

Time difference of 791 days

```
vax$as_of_date[nrow(vax)] - vax$as_of_date[1]
```

Time difference of 784 days

Q9. How many days have passed since the last update of the dataset? 7 days

```
today()-ymd(vax$as_of_date[nrow(vax)])
```

Time difference of 7 days

Q10. How many unique dates are in the dataset (i.e. how many different dates are detailed)?
113 dates

```
length(unique(vax$as_of_date))
```

```
[1] 113
```

Working with ZIP codes

ZIP codes are also rather annoying things to work with as they are numeric but not in the conventional sense of doing math.

Just like dates we have special packages to help us work with ZIP codes.

```
library(zipcodeR)
```

```
geocode_zip('92037')
```

```
# A tibble: 1 x 3
  zipcode lat lng
  <chr>   <dbl> <dbl>
1 92037   32.8 -117.
```

```
zip_distance('92037','95148')
```

```
zipcode_a zipcode_b distance
1      92037      95148    405.6
```

```
reverse_zipcode(c('92037', "95148")) )
```

```
# A tibble: 2 x 24
  zipcode zipcode_~1 major_~2 post_~3 common_c~4 county state lat lng timez~5
  <chr>   <chr>       <chr>   <chr>       <blob> <chr> <chr> <dbl> <dbl> <chr>
1 92037   Standard    La Jol~ La Jol~ <raw 20 B> San D~ CA    32.8 -117. Pacific
2 95148   Standard    San Jo~ San Jo~ <raw 20 B> Santa~ CA    37.3 -122. Pacific
# ... with 14 more variables: radius_in_miles <dbl>, area_code_list <blob>,
#   population <int>, population_density <dbl>, land_area_in_sqmi <dbl>,
#   water_area_in_sqmi <dbl>, housing_units <int>,
#   occupied_housing_units <int>, median_home_value <int>,
#   median_household_income <int>, bounds_west <dbl>, bounds_east <dbl>,
#   bounds_north <dbl>, bounds_south <dbl>, and abbreviated variable names
#   1: zipcode_type, 2: major_city, 3: post_office_city, ...
```

Focus on the San Diego area

Let's now focus on the San Diego county

```
# Subset to San Diego county only areas
sd <- vax[ vax$county == "San Diego" , ]
```

```
library(dplyr)
```

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

filter, lag

The following objects are masked from 'package:base':

intersect, setdiff, setequal, union

```
sd <- filter(vax, county == "San Diego")
```

```
nrow(sd)
```

```
[1] 12091
```

```
sd.10 <- filter(vax, county == "San Diego" &  
  age5_plus_population > 10000)
```

```
head(sd.10)
```

	as_of_date	zip_code_tabulation_area	local_health_jurisdiction	county
1	2021-01-05	92040	San Diego	San Diego
2	2021-01-05	92154	San Diego	San Diego
3	2021-01-05	92122	San Diego	San Diego
4	2021-01-05	92120	San Diego	San Diego
5	2021-01-05	92115	San Diego	San Diego
6	2021-01-05	92114	San Diego	San Diego
	vaccine_equity_metric_quartile		vem_source	
1		3	Healthy Places Index Score	
2		2	Healthy Places Index Score	
3		4	Healthy Places Index Score	
4		4	Healthy Places Index Score	
5		2	Healthy Places Index Score	
6		2	Healthy Places Index Score	
	age12_plus_population	age5_plus_population	tot_population	
1	39405.0	42833	46306	
2	76365.2	82971	88979	
3	44091.1	45951	48071	
4	26372.9	28414	30550	
5	56152.4	60409	64343	

6	59050.7	64945	68851
	persons_fully_vaccinated	persons_partially_vaccinated	
1	14		585
2	16		1397
3	19		1249
4	25		906
5	28		874
6	12		1213
	percent_of_population_fully_vaccinated		
1		0.000302	
2		0.000180	
3		0.000395	
4		0.000818	
5		0.000435	
6		0.000174	
	percent_of_population_partially_vaccinated		
1		0.012633	
2		0.015700	
3		0.025982	
4		0.029656	
5		0.013583	
6		0.017618	
	percent_of_population_with_1_plus_dose	booster_recip_count	
1		0.012935	NA
2		0.015880	NA
3		0.026377	NA
4		0.030474	NA
5		0.014018	NA
6		0.017792	NA
	bivalent_dose_recip_count	eligible_recipient_count	
1	NA		14
2	NA		16
3	NA		19
4	NA		25
5	NA		28
6	NA		12

redacted

1 Information redacted in accordance with CA state privacy requirements

2 Information redacted in accordance with CA state privacy requirements

3 Information redacted in accordance with CA state privacy requirements

4 Information redacted in accordance with CA state privacy requirements

5 Information redacted in accordance with CA state privacy requirements

6 Information redacted in accordance with CA state privacy requirements

Q11. How many distinct zip codes are listed for San Diego County? 107 zipcodes

```
length(unique(sd$zip_code_tabulation_area))
```

```
[1] 107
```

Q12. What San Diego County Zip code area has the largest 12 + Population in this dataset?
92154

```
sd$zip_code_tabulation_area[which.max(sd$age12_plus_population)]
```

```
[1] 92154
```

```
x <-sd %>%  
  filter(as_of_date == "2023-02-28")
```

Q13. What is the overall average “Percent of Population Fully Vaccinated” value for all San Diego “County” as of “2022-02-28”? 74%


```
sd.today <- filter(sd, as_of_date == "2023-02-28")  
  
mean(sd.today$percent_of_population_fully_vaccinated, na.rm = T)
```

```
[1] 0.7400878
```

Q14. Using either ggplot or base R graphics make a summary figure that shows the distribution of Percent of Population Fully Vaccinated values as of “2022-02-28”?

```
library(ggplot2)  
  
ggplot(sd.today, aes(x=sd.today$percent_of_population_fully_vaccinated, )) +  
  geom_histogram(bins = 10, na.rm=T) +  
  labs(x= "Percent of Population Fully Vaccinated in a Zip C0de Area", y="Count (Zip code  
  ggtitle("Histogram of Vaccination Rates Across San Diego County"))
```

Warning: Use of `sd.today\$percent_of_population_fully_vaccinated` is discouraged.
i Use `percent_of_population_fully_vaccinated` instead.




```
ucsd <- filter(sd, zip_code_tabulation_area=="92037")
ucsd[1,]$age5_plus_population
```

```
[1] 36144
```

Q15. Using ggplot make a graph of the vaccination rate time course for the 92037 ZIP code area:

```
ggplot(ucsd) +
  aes(x=as_of_date,
      y=percent_of_population_fully_vaccinated) +
  geom_point() +
  geom_line(group=1) +
  ylim(c(0,1)) +
  labs(title = "Vaccination rate for La Jolla CA 92037", x= "Date", y="Percent Vaccinated")
```

Vaccination rate for La Jolla CA 92037



```
# Subset to all CA areas with a population as large as 92037
vax.36 <- filter(vax, age5_plus_population > 36144 &
  as_of_date == "2022-11-15")

head(vax.36)
```

	as_of_date	zip_code_tabulation_area	local_health_jurisdiction	county
1	2022-11-15	90022	Los Angeles	Los Angeles
2	2022-11-15	92346	San Bernardino	San Bernardino
3	2022-11-15	92231	Imperial	Imperial
4	2022-11-15	95404	Sonoma	Sonoma
5	2022-11-15	92253	Riverside	Riverside
6	2022-11-15	92345	San Bernardino	San Bernardino

	vaccine_equity_metric_quartile	vem_source
1	1	Healthy Places Index Score
2	2	Healthy Places Index Score
3	1	Healthy Places Index Score
4	3	Healthy Places Index Score
5	3	Healthy Places Index Score
6	1	Healthy Places Index Score

	age12_plus_population	age5_plus_population	tot_population
1	55192.3	62369	67014


2	52408.8	58874	63857
3	32448.6	36867	40064
4	35138.9	38436	40497
5	35435.2	38922	40929
6	66047.5	75539	82110
	persons_fully_vaccinated	persons_partially_vaccinated	
1	49419	5537	
2	34564	3266	
3	70739	38805	
4	30573	2342	
5	27448	3453	
6	41672	4283	
	percent_of_population_fully_vaccinated		
1	0.737443		
2	0.541272		
3	1.000000		
4	0.754945		
5	0.670625		
6	0.507514		
	percent_of_population_partially_vaccinated		
1	0.082625		
2	0.051146		
3	0.968575		
4	0.057831		
5	0.084366		
6	0.052162		
	percent_of_population_with_1_plus_dose	booster_recip_count	
1	0.820068	21942	
2	0.592418	16624	
3	1.000000	26706	
4	0.812776	17977	
5	0.754991	14896	
6	0.559676	16576	
	bivalent_dose_recip_count	eligible_recipient_count	redacted
1	4061	49392	No
2	3818	34542	No
3	2737	70617	No
4	7156	30525	No
5	4804	27434	No
6	3191	41649	No

Q16. Calculate the mean “Percent of Population Fully Vaccinated” for ZIP code areas with

a population as large as 92037 (La Jolla) as_of_date “2022-11-15”. Add this as a straight horizontal line to your plot from above with the `geom_hline()` function?

```
average_percent <- mean(vax.36$percent_of_population_fully_vaccinated)
```

```
ggplot(ucsd) +  
  aes(x=as_of_date,  
       y=percent_of_population_fully_vaccinated) +  
  geom_point() +  
  geom_line(group=1) +  
  ylim(c(0,1)) +  
  labs(title = "Vaccination rate for La Jolla CA 92037", x= "Date", y="Percent Vaccinated")  
  geom_hline(yintercept = average_percent,color = "red", linetype=10)
```



Q17. What is the 6 number summary (Min, 1st Qu., Median, Mean, 3rd Qu., and Max) of the “Percent of Population Fully Vaccinated” values for ZIP code areas with a population as large as 92037 (La Jolla) as_of_date “2022-11-15”?

```
summary(vax.36$percent_of_population_fully_vaccinated)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
------	---------	--------	------	---------	------


0.3784 0.6444 0.7162 0.7191 0.7882 1.0000

Q18. Using ggplot generate a histogram of this data.

```
ggplot(vax.36) +  
  aes(x=percent_of_population_fully_vaccinated) +  
  geom_histogram() +  
  labs(x= "Percent Vaccinated", y="Count") +  
  xlim(0,1)
```

`stat_bin()` using `bins = 30`. Pick better value with `binwidth`.

Warning: Removed 2 rows containing missing values (`geom_bar()`).



Q19. Is the 92109 and 92040 ZIP code areas above or below the average value you calculated for all these above? Both averages are lower

```
x <- filter(vax.36, zip_code_tabulation_area %in% c("92109", "92040"))  
x$percent_of_population_fully_vaccinated
```

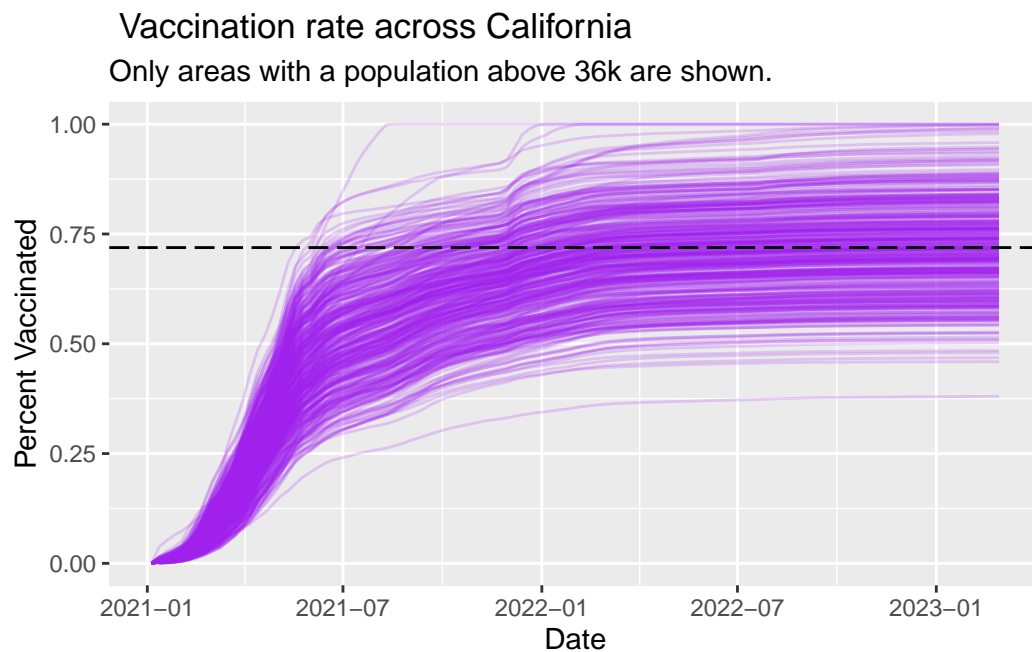
[1] 0.548849 0.692874

```
vax %>% filter(as_of_date == "2022-11-15") %>%
  filter(zip_code_tabulation_area=="92040") %>%
  select(percent_of_population_fully_vaccinated)
```

```
percent_of_population_fully_vaccinated
1                                0.548849
```

```
vax.36.all <- filter(vax,age5_plus_population > 36144)
```

```
ggplot(vax.36.all) +
  aes(as_of_date,
      percent_of_population_fully_vaccinated,
      group=zip_code_tabulation_area) +
  geom_line(alpha=0.2, color="purple", na.rm = T) +
  ylim(0,1) +
  labs(x="Date", y="Percent Vaccinated",
       title= " Vaccination rate across California",
       subtitle="Only areas with a population above 36k are shown.") +
  geom_hline(yintercept = average_percent, linetype=5)
```



Q21. How do you feel about traveling for Thanksgiving Break and meeting for in-person class afterwards? Would be tiring, but I would be okay with it!